

The background is a solid light blue color. Overlaid on this background is a faint, semi-transparent image. On the left side, there is a large stethoscope. In the center, there is a silhouette of a family consisting of two adults and two children, all holding hands in a circle.

MEDICAL INSURANCE

GROUP 14 | BHAVANS VIVEKANANDA COLLEGE
ANSHUMAAN YADAV
SAMYUKTHA SATULURI
ABHINAV SHOURY
G. NAMRATH
UJJWAL KUMAR



ABSTRACT

The project focuses on building a predictive model to forecast insurance premiums based on customer profiles and risk factors, aiming to enhance accuracy in pricing policies and minimize loss ratios.

The study explores a range of machine learning algorithms, including Logistic Regression, KNN, Decision Trees, to predict premium amounts based on factors such as age, vehicle type, and claim history

OBJECTIVE

To identify the most suitable machine learning model for predicting insurance charges in the medical industry, with the goal of optimizing premium plans and improving pricing accuracy.

CONTENT

- Introduction
- Data Preprocessing
- Exploratory Data Analysis
- Data Modelling And Evaluation
- Summary





INTRODUCTION

The Medical Insurance Charges Dataset is an essential resource for analysing healthcare costs and identifying the factors that impact insurance premiums. Key attributes such as age, BMI, smoking habits, and region play a crucial role in determining the cost of medical coverage.

Technological advancements: The use of machine learning and data analytics has revolutionized the prediction of insurance premiums, enabling more precise forecasting and providing valuable insights for optimizing policy pricing and personalized healthcare plans.

DATA PREPROCESSING



DATA

DATASET – The dataset consists of 7 variables and 1339 records.

VARIABLES –

Categorical	Numerical
Region	Age
Sex	BMI
Smoker	Children
	Charges

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

DATA CLEANING

1. Data Overview:

- Inspected the dataset structure (`df.info()`, `df.describe()`, `df.shape`) and previewed the first few rows.

2. Missing Values:

- Checked for missing data using `df.isnull().sum()` (no missing values found).

3. Categorical and Numerical Analysis:

- Used `value_counts()` to check for outliers and anomalies in key columns like age, sex, smoker status, and charges.

DATA CLEANING

4. Encoding Categorical Variables:

- Applied one-hot encoding (`pd.get_dummies`) to convert categorical variables into numeric form for modeling.


5. Multicollinearity Check:

- Calculated Variance Inflation Factor (VIF) to check for multicollinearity between numeric features.

EXPLORATORY DATA ANALYSIS

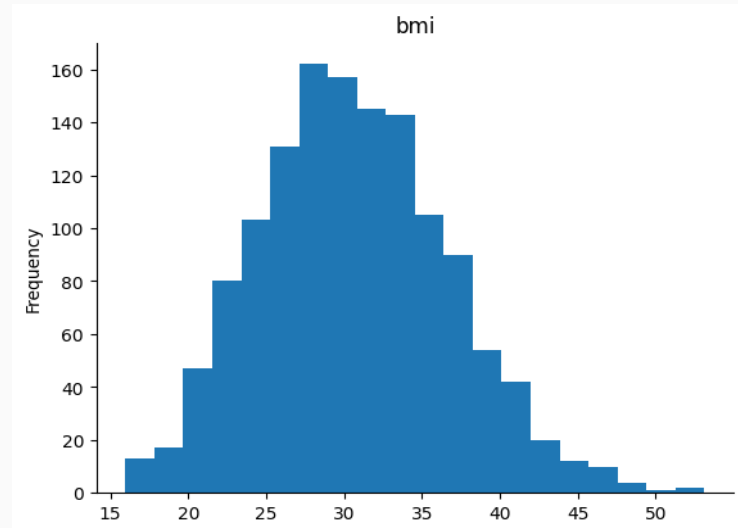

An isometric illustration of a medical data analysis environment. The scene is dominated by a large laptop on the left, displaying a human figure with data points and a line graph. To its right is a tablet showing a heart icon and a waveform. In the foreground, a stethoscope lies on a surface. A tablet in the bottom right corner displays a grid of icons, including a heart and a person. Various medical and data-related icons are scattered throughout, such as a pill bottle, a syringe, and a DNA helix. Several stylized human figures, including doctors in white coats and patients, are depicted interacting with the technology and each other. The entire image has a blue tint.

CORRELATION MATRIX




- There is a strong correlation between –
- Charges vs. Age: Older individuals tend to have higher insurance charges
 - Charges vs. BMI: Individuals with higher BMI face increased insurance charges
 - Charges vs. Smoking Status: Smokers are charged more for insurance.

HISTOGRAM




PIE CHART

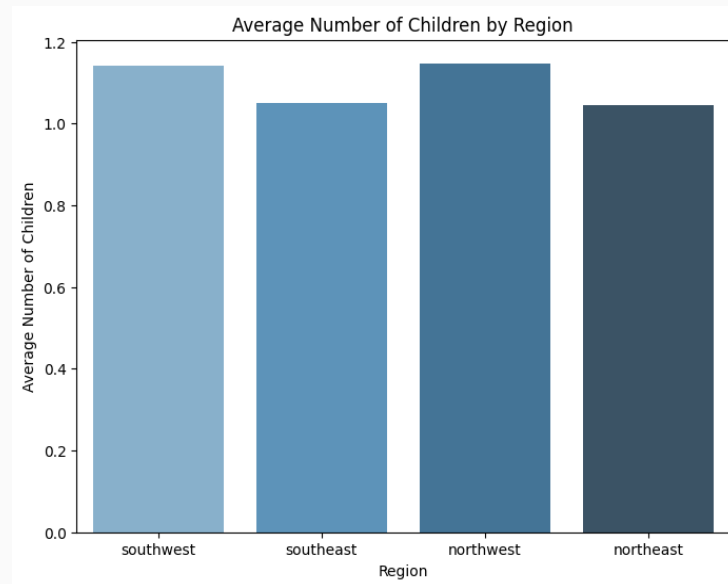

Gender Distribution




Smoker vs Non-Smoker Distribution



BAR PLOT




SCATTER PLOT



MACHINE LEARNING ALGORITHM

(CLASSIFICATION MODEL)



LOGISTIC REGRESSION

Train – test Split Ratio	Accuracy
0.20	91.04%
0.25	89.25%
0.35	89.9%

KNN CLASSIFICATION

Train – test Split Ratio	Accuracy
0.20	79.4%
0.25	77.31%
0.35	76.54%

DECISION TREE CLASSIFIER

Train – test Split Ratio	Accuracy
0.20	90.29%
0.25	90.74%
0.35	91.04%

ALGORITHM COMPARISON

Algorithm	Accuracy
Logistic Regression	91.04%
KNN Classification	77.31%
Decision Tree Classifier	91.04%

SUMMARY

- This study evaluates machine learning models for predicting medical insurance premiums, focusing on identifying the most accurate approach. The models compared include Logistic Regression, K-Nearest Neighbor (KNN), and Decision Tree Classifier.
- Logistic Regression and Decision Tree Classifier both achieved the highest accuracy at 91.04%, making them equally effective in this context.
- KNN Classification had the lowest performance, with an accuracy of 77.31%.
- The results suggest that Logistic Regression and Decision Tree models are equally suitable for predicting insurance premiums, while KNN underperforms in comparison

Thank You!

- ANSHUMAAN YADAV
- SAMYUKTHA SATULURI
- ABHINAV SHOURY
- G NAMRATH
- UJJWAL KUMAR



APPENDIX



Loading the Dataset

```
df=pd.read_csv('/content/Insurance Premium.csv')  
df
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

Null Values

Checking for the data type

```
df.isnull().sum()
```

```
0
age    0
sex    0
bmi    0
children 0
smoker 0
region 0
charges 0
```

```
dtype: int64
```

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
None
```


Calculating Variance Inflation Factor (VIF) for Feature Selection in Regression Analysis

```
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Select only numeric columns (ignoring any non-numeric)
df_numeric = df_dummy.select_dtypes(include=[np.number])

# Create DataFrame to store VIF values
vif_data = pd.DataFrame()
vif_data["Feature"] = df_numeric.columns
vif_data["VIF"] = [variance_inflation_factor(df_numeric.values, i) for i in range(df_numeric.shape[1])]

print(vif_data)
```

	Feature	VIF
0	age	8.098132
1	bmi	8.044400
2	children	1.800015
3	charges	2.473524


Splitting Dataset into Training and Testing Sets for Regression Analysis

```
X = df_dummy.drop(['charges'],axis=1)
y = df_dummy['charges']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.4, random_state=42)
print(X_train.shape,y_train.shape,X_test.shape,y_test.shape)
```


```
(802, 8) (802,) (536, 8) (536,)
```

Plots- HIST

```
df['bmi'].plot(kind='hist', bins=20, title='bmi')  
plt.gca().spines[['top', 'right',]].set_visible(False)
```



```
from matplotlib import pyplot as plt  
df['children'].plot(kind='hist', bins=20, title='children')  
plt.gca().spines[['top', 'right',]].set_visible(False)
```



Logistic Regression

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
```

```
df1 = df_dummy.copy()
```

```
charges_med = df1['charges'].median()
df1['catcharges'] = [1 if x > charges_med else 0 for x in df1['charges']]
```

```
df1['catcharges'].value_counts(normalize = True)
df1 = df1.drop('charges', axis=1)
```

```
x = df1.drop(['catcharges'],axis=1)
Y = df1['catcharges']
x_train, x_test, Y_train, Y_test = train_test_split(x, Y, test_size=.2, random_state=42)
print(x_train.shape,Y_train.shape,x_test.shape,Y_test.shape)
```

```
(1070, 8) (1070,) (268, 8) (268,)
```

```
model3=LogisticRegression()
model3.fit(x_train,Y_train)
```

```
model3.coef_
```

```
array([[ 0.14663096,  0.03259212,  0.05972964, -0.23076158,  6.1251742 ,
        -0.32717959, -0.55202742, -0.65595449]])
```

```
Y_pred = model3.predict(x_test)
Y_pred
```

```
array([1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0,
       0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1,
       0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0,
       1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1,
       0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1,
       0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0,
       0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1,
       0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0,
       0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1,
       0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1,
       1. 0. 1. 1])
```

```
# binary, multiclass
```

```
from sklearn.metrics import classification_report
classification_rep = classification_report(Y_test, Y_pred)
print(classification_rep)
```

	precision	recall	f1-score	support
0	0.94	0.90	0.92	146
1	0.88	0.93	0.90	122
accuracy			0.91	268
macro avg	0.91	0.91	0.91	268
weighted avg	0.91	0.91	0.91	268

K-Nearest Neighbor

```
[ ] df_dummy.shape
```


```
⇒ (1338, 9)
```

```
[ ] print(1338 ** 0.5)
```

```
⇒ 36.578682316343766
```

```
[ ] from sklearn.neighbors import KNeighborsClassifier  
model2=KNeighborsClassifier(n_neighbors=37)
```

```
▶ model2.fit(x_train, Y_train)
```

```
⇒  KNeighborsClassifier ⓘ ?  
KNeighborsClassifier(n_neighbors=37)
```

```
Y_pred = model2.predict(x_test)
```

```
Y_pred
```

```
array([0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0,  
       0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1,  
       0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0,  
       0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1,  
       0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1,  
       0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0,  
       0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0,  
       0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1,  
       0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1,  
       0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0,  
       0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0,  
       0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1,  
       1, 0, 1, 1])
```

```
[ ] accuracy_score(Y_test,Y_pred)
```

```
⇒ 0.7910447761194029
```

```
[ ] cm2 = confusion_matrix(Y_test,Y_pred)  
cm2
```

```
⇒ array([[135, 11],  
         [ 45, 77]])
```

```
[ ] metrics.mean_absolute_error(Y_test,Y_pred)
```

```
⇒ 0.208955223880597
```

Decision Tree

```
[ ] feature_cols = ['age', 'bmi', 'children', 'sex_male', 'smoker_yes', 'region_northwest', 'region_southeast', 'region_southwest']
s = df1[feature_cols] # Features
r = df1['catcharges']
s_train, s_test, r_train, r_test = train_test_split(s, r, test_size=0.2, random_state=1) # 70% training and 30% test
```

```
[ ] from sklearn.tree import DecisionTreeClassifier
```

```
clf = DecisionTreeClassifier()
```

```
clf = clf.fit(s_train, r_train)
```

```
▶ r_pred = clf.predict(s_test)
print("Accuracy:", metrics.accuracy_score(r_test, r_pred))
```

```
↔ Accuracy: 0.9029850746268657
```

```
[ ] clf = DecisionTreeClassifier(criterion="entropy", max_depth=3)
```

```
clf = clf.fit(s_train, r_train)
```

```
r_pred = clf.predict(s_test)
```

```
print("Accuracy:", metrics.accuracy_score(r_test, r_pred))
```

```
↔ Accuracy: 0.917910447761194
```

```
[ ] clf = DecisionTreeClassifier(criterion="gini", max_depth=2)

clf = clf.fit(s_train, r_train)

y_pred = clf.predict(s_test)

print("Accuracy:", metrics.accuracy_score(r_test, r_pred))
```

```
↔ Accuracy: 0.917910447761194
```

```
[ ] clf = DecisionTreeClassifier(criterion="gini", max_depth=2)

clf = clf.fit(s_train, r_train)

y_pred = clf.predict(s_test)

print("Accuracy:", metrics.accuracy_score(r_test, r_pred))
```

```
↔ Accuracy: 0.917910447761194
```

```
▶ clf = DecisionTreeClassifier(criterion="gini", max_depth=3)

clf = clf.fit(s_train, r_train)

y_pred = clf.predict(s_test)

print("Accuracy:", metrics.accuracy_score(r_test, r_pred))
```

```
↔ Accuracy: 0.917910447761194
```

Random Forest

```
[ ] from sklearn.datasets import load_iris
    from sklearn.tree import DecisionTreeClassifier
    from sklearn.ensemble import RandomForestClassifier
    from sklearn.model_selection import train_test_split
    from sklearn.metrics import classification_report, confusion_matrix
    from sklearn.tree import plot_tree

    from matplotlib import pyplot as plt
```

```
[ ] A = df1[['age', 'bmi', 'children', 'sex_male', 'smoker_yes', 'region_northwest', 'region_southeast', 'region_southwest']]
    B = df1['catcharges']
```

```
[ ] A_train, A_test, B_train, B_test = train_test_split(A, B, test_size=0.2)
```

```
▶ print(A_train.shape)
   print(B_train.shape)
   print(A_test.shape)
   print(B_test.shape)
```

```
⇒ (1070, 8)
   (1070,)
   (268, 8)
   (268,)
```

```
[ ] rf = RandomForestClassifier()
```

```
[ ] rf.fit(A_train, B_train)
```



```
▼ RandomForestClassifier ⓘ ?
RandomForestClassifier()
```

```
[ ] B_pred = rf.predict(A_test)
```



```
print(classification_report(B_test, B_pred))
print(confusion_matrix(B_test, B_pred))
```



	precision	recall	f1-score	support
0	0.91	0.95	0.93	135
1	0.94	0.90	0.92	133
accuracy			0.93	268
macro avg	0.93	0.93	0.93	268
weighted avg	0.93	0.93	0.93	268

```
[[128  7]
 [ 13 120]]
```